

COMP 474/6741 Intelligent Systems (Winter 2024)

Worksheet #6: Intelligent Agents

Task 1. What kind of question would you expect a *Concordia Chatbot* to be able to handle?

1. For *new students* (not currently registered at Concordia):

- [How do I enroll in a class?](#)

-

-

2. For *current students* (already enrolled at Concordia):

-

-

-

Task 2. Write a *regular expression* that matches different variations of naming <https://www.wikidata.org/wiki/Q326342>: “Concordia”, “Concordia U.”, “CU”, “Concordia University”, “Université Concordia”, ...:

[https://www.wikidata.org/wiki/Q326342|Concordia\(\\sU\\.?\)|\(\\s\)?University\)?|CU|Universit\(é|e\)\\sConcordia](https://www.wikidata.org/wiki/Q326342|Concordia(\\sU\\.?)|(\\s)?University)?|CU|Universit(é|e)\\sConcordia)

Test it at <https://regex101.com/>

\\s = whitespace

\\.?: This means that the dot . is optional.

Task 3. Suppose we want to use an existing set of 1000 questions for training a ML classifier. If we use tf-idf vectors to represent each question, how many dimensions will the vectors have (make a rough estimate)?

10 unique words on average

1000 * 10 = 10,000 unique word set

Task 4. Ok, here is an (extremely simplified) idea of creating 2D feature vectors out of a natural language question: The first dimension *a* encodes the first occurrence of a question word (see table below) and the second dimension *b* the number of Capital Letters in the sentence:

Contains?	Value	#	Question	a	b	Class
Who	1	1	Where is Concordia?	3	1	Location
What	2	2	Who was Steve Jobs?	1	2	Definition
Where	3	3	What city is McGill in?	2	2	Location
(none)	0	4	What is NLP?	2	1	Definition



Task 5. Using the online parser at <https://corenlp.run/>, create a *parse tree* for the sentence *What is McGill?*. Note that you can now extract the *subject* of the sentence, e.g., to plug it into a SPARQL query.

Results: SUBJECT = McGill

Task 6. Now apply the kNN classification algorithm on the new question below to classify its type, according to the training data from Task 4. Use $k = 3$ and the Euclidian distance $d(\vec{p}, \vec{q}) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$:

#	Question	a	b	d-Q1	d-Q2	d-Q3	d-Q4	Class?
5	What is McGill?	2	2	1	1	1	0	Definition


You can now match the new question with a corresponding SPARQL template to obtain a query for your knowledge graph, filling in variables with the values extracted from the question.

$$d_{Q1} = \sqrt{(2-3)^2 + (2-1)^2} = 1$$

The nearest neighbors for $k = 3$ are:

(2, 1) - Definition
(3, 1) - Location
(2, 2) - Definition

Majority voting: The majority class among the nearest neighbors is "Definition".

-  **Task 7.** Now define a SPARQL template that can obtain information about a *person* from DBpedia. To keep it simple, for now we assume that the name extracted via NLP from the question is identical to the full name stored in the (English) `label` field for the subject (e.g., “Steve Jobs”).

```

PREFIX dbr: <http://dbpedia.org/resource/>
SELECT . . .
WHERE {
    . . .
    PREFIX dbr: <http://dbpedia.org/resource/>
    SELECT ?property ?value
    WHERE {
        dbr:Steve_Jobs ?property ?value .
    }
}

```

- Task 8.** Create a *competency question* and a corresponding SPARQL query for our FOCU university example:

"What are the courses offered by FOCU University?"


Testing query for this question:

```


SELECT . . .
WHERE {
    . . .
    PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
    PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
    PREFIX dbo: <http://dbpedia.org/ontology/>
    PREFIX foaf: <http://xmlns.com/foaf/0.1/>

    SELECT ?course
    WHERE {
        ?university rdf:type dbo:University ;
            rdfs:label "FOCU University"@en ;
            dbo:offers ?course_resource .
        ?course_resource rdfs:label ?course .
    }
}

```

-  **Task 9.** An early, well-known commercial service for semantic annotation of textual (mostly news) documents was Thompson Reuter’s *OpenCalais*, which has since been spun out and re-branded as *LSEG Data & Analytics PermID* (formerly Refinitif Intelligent Tagging): Try out the online demo at <https://permid.org/tagging> on a document, for example the first part of the Wikipedia article on Concordia. Look at the entities that were detected and go to the “RDF view”: what ID is given to Concordia in this knowledge graph?

Hint: There is another tool at the top of the page, *Entity Search*, where you can cross-check your entities.

-  **Task 10.** Go to the DBpedia *Spotlight* online demo at <https://demo.dbpedia-spotlight.org>. Try analyzing a test document with some ambiguities, e.g., “*Paris Hilton went to the Hilton in Paris.*” Inspect the entities that were linked to DBpedia. Are they correct? no, hovering over it gives:
http://ca.dbpedia.org/resource/Paris_Hilton

Task 11. For the questions in Task 1 above, which of the chatbot techniques covered so far would be able to answer them?